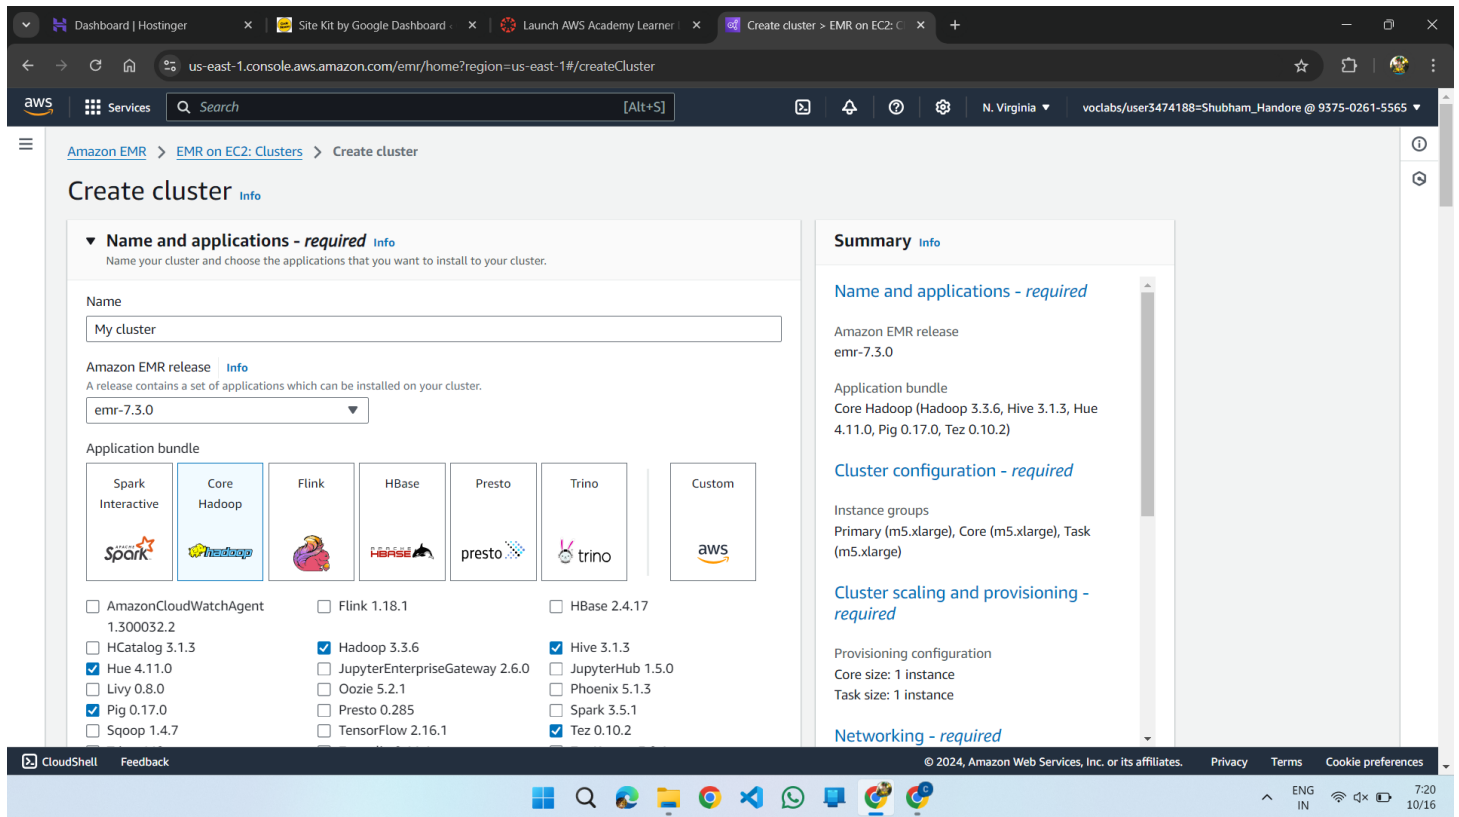
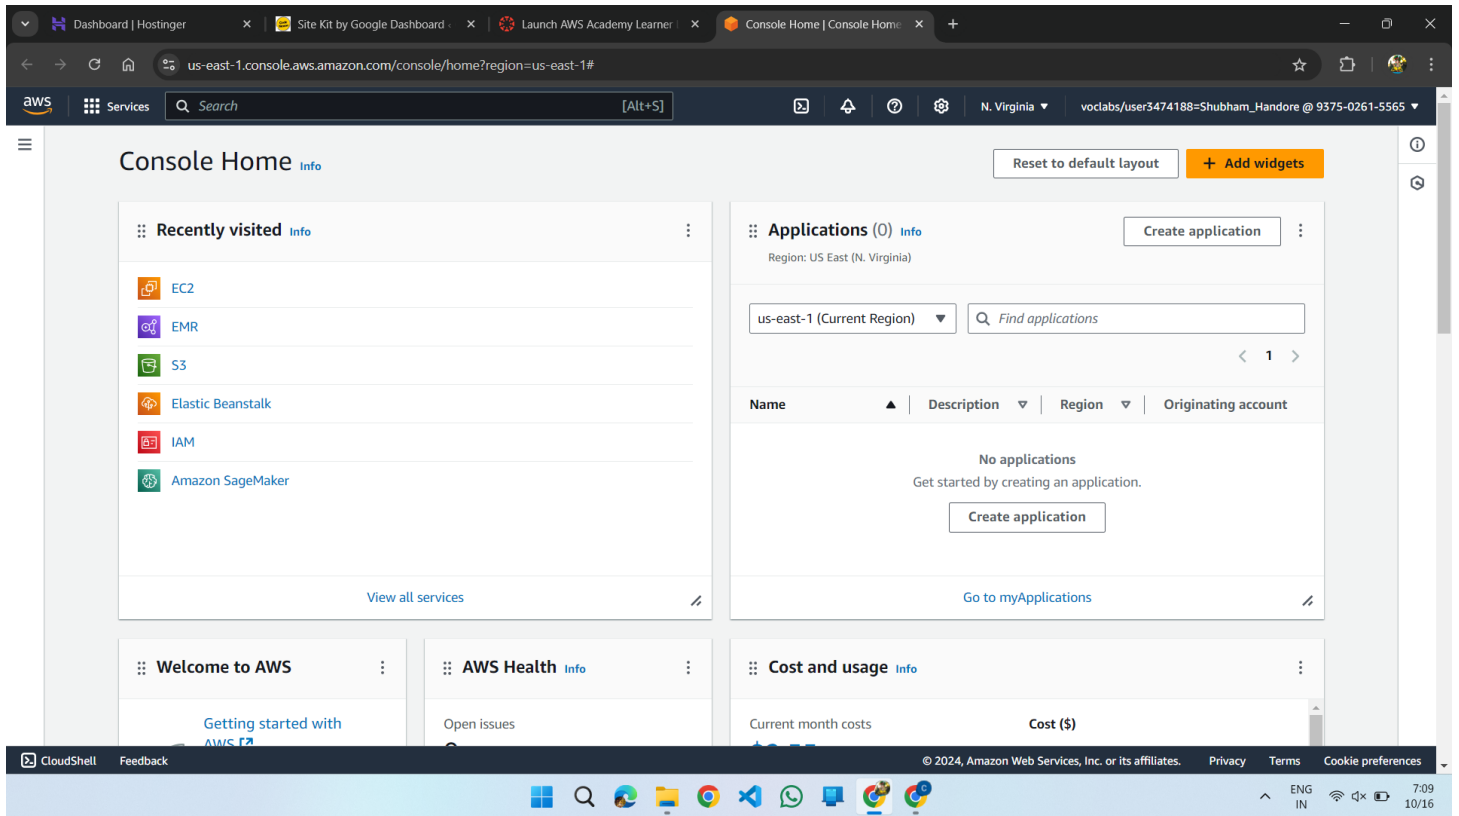
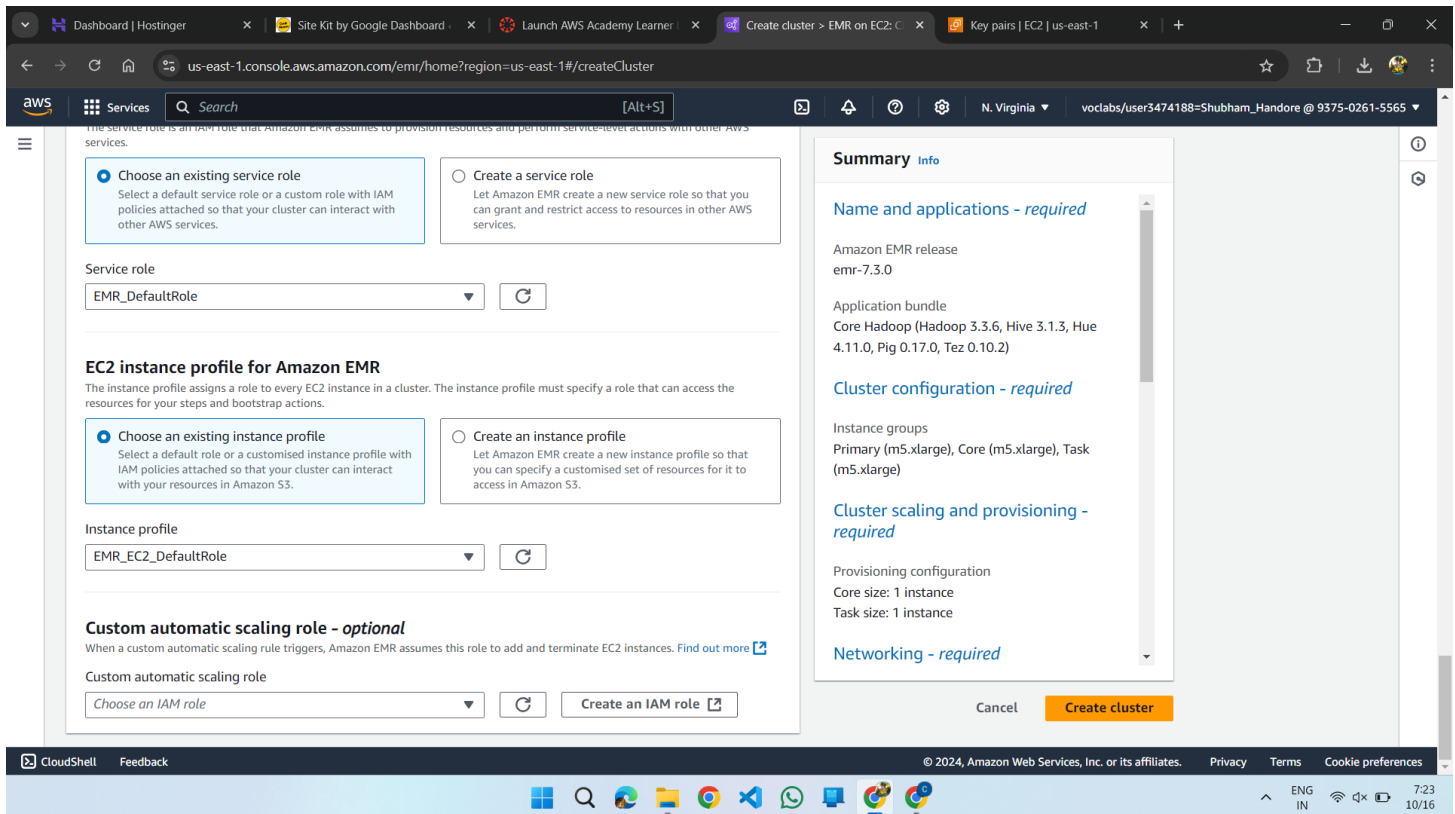
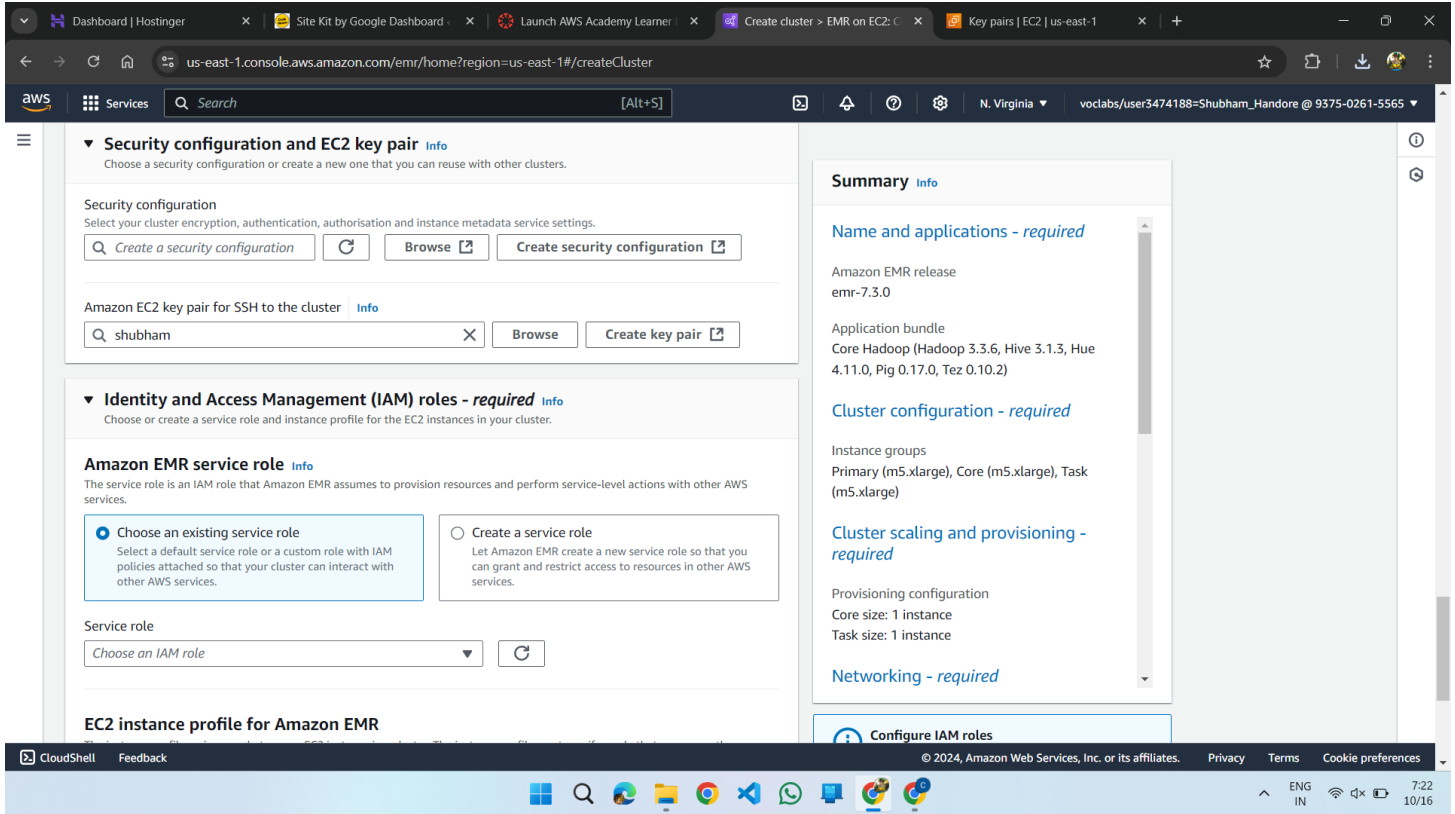


Cloud Computing Assignment 10

Installation and Configuration of cloud Hadoop and demonstrate simple query





Dashboard | HostingerSite Kit by Google DashboardLaunch AWS Academy LearnerProperties > My cluster > EMRKey pairs | EC2 | us-east-1

us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/clusterDetails/j-YEAMIUNJST74

awsServicesSearch[Alt+S]

N. Virginiavoclabs/user3474188=Shubham_Handore @ 9375-0261-5565

Your cluster 'My cluster' has been successfully created.

Amazon EMR> EMR on EC2: Clusters> My cluster

My cluster

Updated less than a minute ago

Terminate

Clone in AWS CLI

Clone

▼ Summary

Cluster info

Cluster ID
j-YEAMIUNJST74

Cluster configuration
Instance groups

Capacity
1 Primary 1 Core 1 Task

Applications

Amazon EMR version
emr-7.3.0

Installed applications
Hadoop 3.3.6, Hive 3.1.3, Hue 4.11.0, Pig 0.17.0, Tez 0.10.2

Cluster management

Log destination in Amazon S3
aws-logs-937502615565-us-east-1/elasticmapreduce

Primary node public DNS
-

Status and time

Status
Starting

Creation time
16 October 2024 19:23 (UTC+05:30)

Elapsed time
1 second

Properties

Bootstrap actions

Instances (hardware)

Steps

Applications

Configurations

Monitoring

Events

Tags (0)

Cluster logs Info

Archive log files to Amazon S3
Turned on

Amazon S3 location
s3://aws-logs-937502615565-us-east-

Turn on encryption for logs
Turned off

Cluster termination and node replacement Info

Termination option
Automatically terminate the cluster after idle time

Termination protection

Idle time
1 hour

Unhealthy node replacement

CloudShellFeedback© 2024, Amazon Web Services, Inc. or its affiliates. PrivacyTermsCookie preferences

Dashboard | HostingerSite Kit by Google DashboardLaunch AWS Academy LearnerInstances (hardware) > My clusterKey pairs | EC2 | us-east-1

us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/clusterDetails/j-YEAMIUNJST74

awsServicesSearch[Alt+S]

N. Virginiavoclabs/user3474188=Shubham_Handore @ 9375-0261-5565

1 Primary 1 Core 1 Task18 seconds

Properties

Bootstrap actions

Instances (hardware)

Steps

Applications

Configurations

Monitoring

Events

Tags (0)

Instance group settings Info

Edit cluster scaling option

Cluster scaling option
Manually set cluster size

Core

Name and maximum core nodes in the cluster
Core 1 instances

Task

Name and maximum task nodes in the cluster
Task - 1 1 instances

Instance groups (3) Info

With the instance groups configuration, each node type consists of the same instance type and the same purchasing option for instances: on-demand or spot.

Find instances by statusFind resources by ID or type; or search for text within loaded results

Type and name

ID

Status
Last state change reason

Instances

Purchasing option and p...

EBS size (GiB)

EC2

Primary

ig-1N8UUUP320O39

Provisioning

0 (1 requested)

On-demand

-

-

Core

ig-SJ1NOV6E71RO

Provisioning

0 (1 requested)

On-demand

-

-

Task (Task - 1)

ig-27E4DQW8IKDS8

Provisioning

0 (1 requested)

On-demand

-

-

CloudShellFeedback© 2024, Amazon Web Services, Inc. or its affiliates. PrivacyTermsCookie preferences

Dashboard | HostingerSite Kit by Google DashboardLaunch AWS Academy LearnerProperties > My cluster > EMRKey pairs | EC2 | us-east-1

us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/clusterDetails/j-YEAMIUNJST74

awsServicesSearch[Alt+S]

N. Virginiavoclabs/user3474188=Shubham_Handore @ 9375-0261-5565

Unhealthy node replacement
On

Network and securityInfo

Network

Virtual Private Cloud (VPC)
vpc-06b72cb8fe3aaabb1

Subnet ID and Availability Zone (AZ)
subnet-00970aad8c04b4b0e us-east-1d

EC2 security groups (firewall)

Primary node
EMR-managed security group
sg-00d3dbefadd8e4f62

Additional security groups
-

Core and task nodes
EMR-managed security group
sg-05990083e4f85131a

Additional security groups
-

Security configuration

Security configuration
none

EC2 key pair
shubham

Permissions

Service role for Amazon EMR
EMR_DefaultRole

EC2 instance profile
EMR_EC2_DefaultRole

Auto-scaling role
Not configured

https://us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#SecurityGroup:grou...

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

ENG IN 7:28 10/16

Dashboard | HostingerSite Kit by Google DashboLaunch AWS Academy LeaProperties > My cluster > EC2 | us-east-1Key pairs | EC2 | us-east-1

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#SecurityGroup:groupId=sg-00d3dbefadd8e4f62

awsServicesSearch[Alt+S]

N. Virginiavoclabs/user3474188=Shubham_Handore @ 9375-0261-5565

EC2 DashboardEC2 Global ViewEventsInstancesInstance TypesLaunch TemplatesSpot RequestsSavings PlansReserved InstancesDedicated HostsCapacity ReservationsImagesAMIsAMI CatalogElastic Block StoreVolumesSnapshotsLifecycle ManagerNetwork & SecurityCloudShellFeedback

Owner
937502615565

Inbound rules count
7 Permission entries

Outbound rules count
1 Permission entry

Inbound rulesOutbound rulesTags

Inbound rules (7)

Search

	Name	Security group rule...	IP version	Type	Protocol	Port range
<input type="checkbox"/>	-	sgr-04c3d7e9d301a10fc	-	All UDP	UDP	0 - 65535
<input type="checkbox"/>	-	sgr-083df835aa3d84865	-	All ICMP - IPv4	ICMP	All
<input type="checkbox"/>	-	sgr-0528be2293862d...	-	All TCP	TCP	0 - 65535
<input type="checkbox"/>	-	sgr-0661951c9fc2cb170	-	Custom TCP	TCP	8443
<input type="checkbox"/>	-	sgr-0db2bf946fb1ead9f	-	All TCP	TCP	0 - 65535
<input type="checkbox"/>	-	sgr-02f3cb43f5d6886a7	-	All UDP	UDP	0 - 65535
<input type="checkbox"/>	-	sgr-00af74fdb73082521	-	All ICMP - IPv4	ICMP	All

Manage tagsEdit inbound rules

CloudShellFeedback

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

ENG IN 7:28 10/16

Dashboard | Hostinger

Site Kit by Google Dashbo

Launch AWS Academy Lea

Properties > My cluster >

ModifyInboundSecurityGr

Key pairs | EC2 | us-east-1

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#ModifyInboundSecurityGroupRules:securityGroupId=sg-00d3dbefadd8e4f62

Services

Search

[Alt+S]

N. Virginia

voclabs/user3474188=Shubham_Handore @ 9375-0261-5565

EC2 > Security Groups > sg-00d3dbefadd8e4f62 - ElasticMapReduce-master > Edit inbound rules

Info

Edit inbound rules

Inbound rules control the incoming traffic that's allowed to reach the instance.

Inbound rules

Info

Security group rule ID	Type	Protocol	Port range	Source	Description - optional	
	Info	Info	Info	Info	Info	
sg-04c3d7e9d301a10fc	All UDP	UDP	0 - 65535	Custom	<div>Q</div>	<div>Delete</div>
sg-083df835aa3d84865	All ICMP - IPv4	ICMP	All	Custom	<div>Q</div> <div>sg-05990083e4f85131a</div>	<div>Delete</div>
sg-0528be2293862d0eb	All TCP	TCP	0 - 65535	Custom	<div>Q</div> <div>sg-00d3dbefadd8e4f62</div>	<div>Delete</div>
sg-0661951c9fc2cb170	Custom TCP	TCP	8443	Custom	<div>Q</div> <div>sg-00d3dbefadd8e4f62</div>	<div>Delete</div>
sg-0db2bf946fb1ead9f	All TCP	TCP	0 - 65535	Custom	<div>Q</div> <div>pl-f8bd5e91</div>	<div>Delete</div>

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

ENG

IN

7:29

10/16

Dashboard | Hostinger

Site Kit by Google Dashbo

Launch AWS Academy Lea

Properties > My cluster >

ModifyInboundSecurityGr

Key pairs | EC2 | us-east-1

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#ModifyInboundSecurityGroupRules:securityGroupId=sg-00d3dbefadd8e4f62

Services

Search

[Alt+S]

N. Virginia

voclabs/user3474188=Shubham_Handore @ 9375-0261-5565

sg-083df835aa3d84865

All ICMP - IPv4

ICMP

All

Custom

Q

sg-05990083e4f85131a

Delete

sg-0528be2293862d0eb

All TCP

TCP

0 - 65535

Custom

Q

sg-00d3dbefadd8e4f62

Delete

sg-0661951c9fc2cb170

Custom TCP

TCP

8443

Custom

Q

sg-00d3dbefadd8e4f62

Delete

sg-0db2bf946fb1ead9f

All TCP

TCP

0 - 65535

Custom

Q

pl-f8bd5e91

Delete

sg-02f3cb43f5d6886a7

All UDP

UDP

0 - 65535

Custom

Q

sg-05990083e4f85131a

Delete

sg-00af74fdb73082521

All ICMP - IPv4

ICMP

All

Custom

Q

sg-00d3dbefadd8e4f62

sg-05990083e4f85131a

Delete

Add rule

Cancel

Preview changes

Save rules

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

ENG

IN

7:30

10/16

Dashboard | HostingerSite Kit by Google DashboLaunch AWS Academy LeProperties > My cluster >ModifyInboundSecurityGrKey pairs | EC2 | us-east-1

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#ModifyInboundSecurityGroupRules:securityGroupId=sg-00d3dbefadd8e4f62

ServicesSearch[Alt+S]

N. Virginiavoclabs/user3474188=Shubham_Handore @ 9375-0261-5565

sg-0661951c9fc2cb170

Custom TCP

TCP

8443

Custom

sg-00d3dbefadd8e4f62

Q

Delete

sg-0db2bf946fb1ead9f

All TCP

TCP

0 - 65535

Custom

pl-f8bd5e91

Q

Delete

sg-02f3cb43f5d6886a7

All UDP

UDP

0 - 65535

Custom

sg-05990083e4f85131a

Q

Delete

sg-00af74fdb73082521

All ICMP - IPv4

ICMP

All

Custom

sg-00d3dbefadd8e4f62

Q

Delete

-

SSH

TCP

22

Anywh...

sg-05990083e4f85131a

Q 0.0.0.0/0

Delete

0.0.0.0/0

Add rule

Rules with source of 0.0.0.0/0 or ::/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.

Cancel

Preview changes

Save rules

CloudShellFeedback

© 2024, Amazon Web Services, Inc. or its affiliates. PrivacyTermsCookie preferences

ENG IN7:30 10/16

Dashboard | HostingerSite Kit by Google DashboLaunch AWS Academy LeProperties > My cluster >SecurityGroup | EC2 | us-eKey pairs | EC2 | us-east-1

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#SecurityGroup:group-id=sg-00d3dbefadd8e4f62

ServicesSearch[Alt+S]

N. Virginiovoclabs/user3474188=Shubham_Handore @ 9375-0261-5565

EC2 Dashboard

EC2 Global View

Events

Instances

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Capacity Reservations

Images

AMIs

AMI Catalog

Elastic Block Store

Volumes

Snapshots

Lifecycle Manager

Network & Security

Inbound security group rules successfully modified on security group (sg-00d3dbefadd8e4f62 | ElasticMapReduce-master)

Details

EC2 > Security Groups > sg-00d3dbefadd8e4f62 - ElasticMapReduce-master

sg-00d3dbefadd8e4f62 - ElasticMapReduce-master

Actions

Details

Security group name

ElasticMapReduce-master

Security group ID

sg-00d3dbefadd8e4f62

Description

Master group for Elastic MapReduce created on 2024-09-29T02:30:00.524Z

VPC ID

vpc-06b72cb8fe3aaabb1

Owner

937502615565

Inbound rules count

8 Permission entries

Outbound rules count

1 Permission entry

Inbound rules

Outbound rules

Tags

Inbound rules (8)

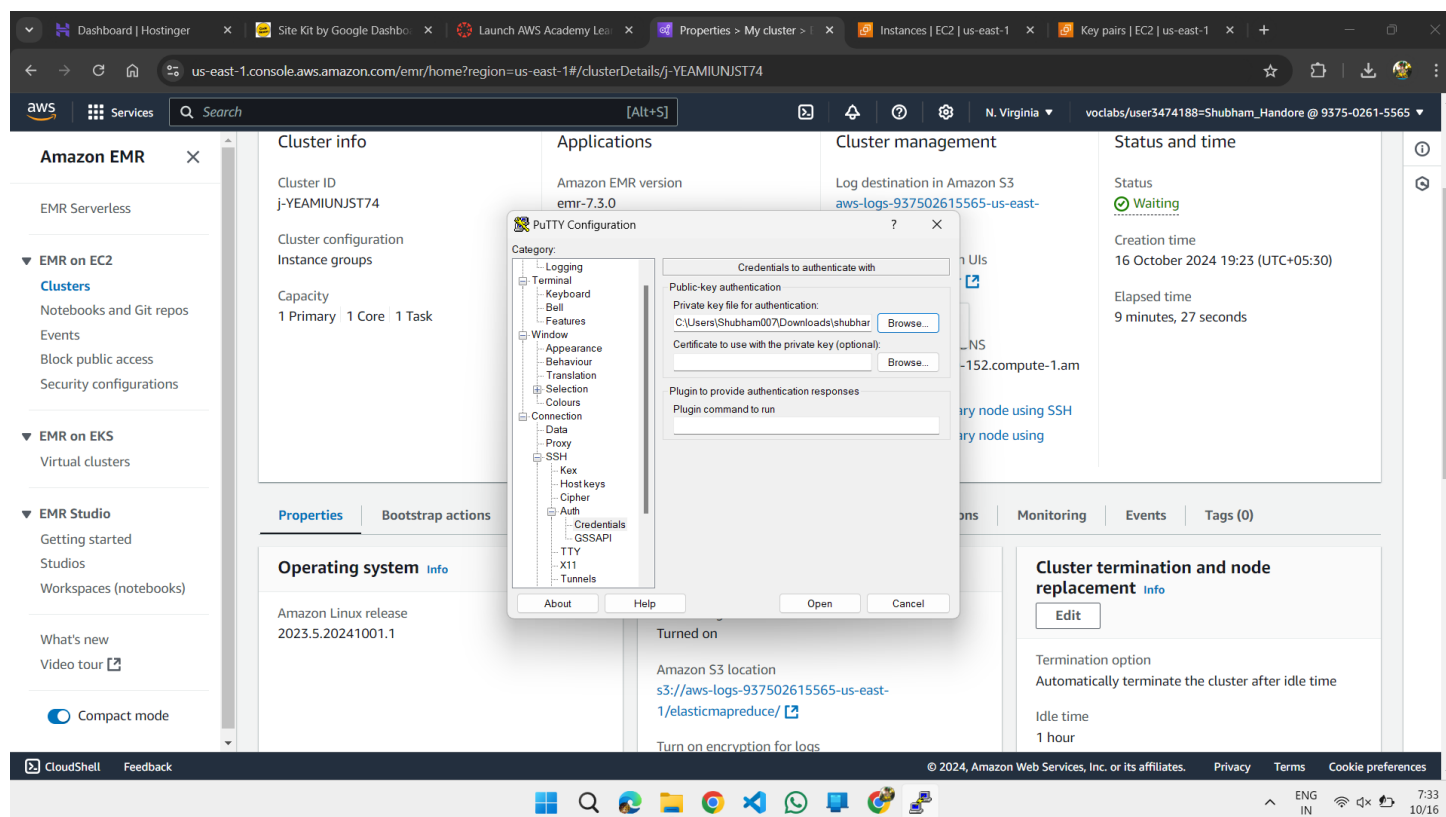
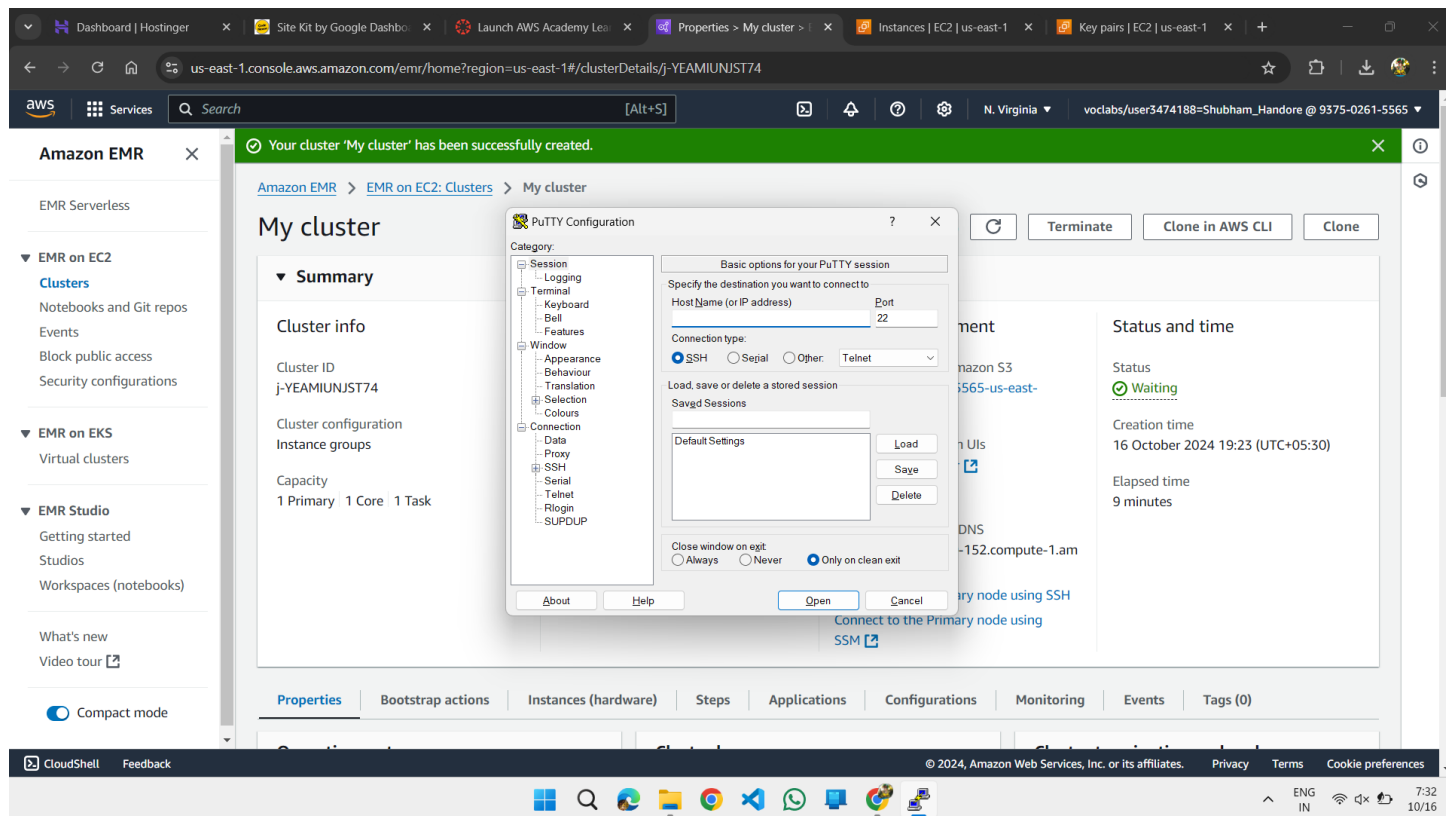
Manage tags

Edit inbound rules

CloudShellFeedback

© 2024, Amazon Web Services, Inc. or its affiliates. PrivacyTermsCookie preferences

ENG IN7:30 10/16




```

at org.apache.hadoop.hive ql.parse.HiveParser.queryStatementExpressionBody(HiveParser.java:39438)
at org.apache.hadoop.hive ql.parse.HiveParser.queryStatementExpression(HiveParser.java:39326)
at org.apache.hadoop.hive ql.parse.HiveParser.execStatement(HiveParser.java:2482)
at org.apache.hadoop.hive ql.parse.HiveParser.statement(HiveParser.java:1435)
at org.apache.hadoop.hive ql.parse.ParseDriver.parse(ParseDriver.java:230)
at org.apache.hadoop.hive ql.parse.ParseUtils.parse(ParseUtils.java:79)
at org.apache.hadoop.hive ql.parse.ParseUtils.parse(ParseUtils.java:72)
at org.apache.hadoop.hive ql.Driver.compile(Driver.java:631)
at org.apache.hadoop.hive ql.Driver.compileInternal(Driver.java:1871)
at org.apache.hadoop.hive ql.Driver.compileAndRespond(Driver.java:1818)
at org.apache.hadoop.hive ql.Driver.compileAndRespond(Driver.java:1813)
at org.apache.hadoop.hive ql.rexec.ReExecDriver.compileAndRespond(ReExecDriver.java:126)
at org.apache.hadoop.hive ql.rexec.ReExecDriver.run(ReExecDriver.java:214)
at org.apache.hadoop.hive cli.CliDriver.processLocalCmd(CliDriver.java:242)
at org.apache.hadoop.hive cli.CliDriver.processCmd(CliDriver.java:191)
at org.apache.hadoop.hive cli.CliDriver.processLine(CliDriver.java:405)
at org.apache.hadoop.hive cli.CliDriver.executeDriver(CliDriver.java:849)
at org.apache.hadoop.hive cli.CliDriver.run(CliDriver.java:787)
at org.apache.hadoop.hive cli.CliDriver.main(CliDriver.java:709)
at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:77)
at java.base/jdk.internal.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.base/java.lang.reflect.Method.invoke(Method.java:569)
at org.apache.hadoop.util.RunJar.run(RunJar.java:328)
at org.apache.hadoop.util.RunJar.main(RunJar.java:241)
FAILED: ParseException line 1:39 cannot recognize input near '1' '',' '01' in constant
hive> insert into student values ('shubham',1,'01'),('nitish',2,'02');
Query ID = root_20241016141022_86393423-f017-4a90-aiad-3dcb8e73b6ad
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1729087094069_0002)

```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0

VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 5.12 s

```

Loading data to table default.student
OK
Time taken: 7.646 seconds
hive> select * from student;
OK
shubham 1      1.0
nitish  2      2.0
Time taken: 0.293 seconds, Fetched: 2 row(s)
hive> █

```

